



Machine Learning Models and Dataset Versioning



\$ whoami

- Open source contributor DVC, CloudCV,
- FOSSASIA OpenTechNights Winner
- ML intern @Neuroplex
- ML and AI enthusiast
- Final year Compute science student @ MEC



Outline

- Why ML is different from software engineering?
- Problems on using git
- Problems with git-LFS
- Current workflow for ML models management
- Use case as application- Cats vs Dogs model
- Conclusion



What is ML?

- It provides an systems the ability to automatically learn and improve from experience without being explicitly programmed.
- (supervised, unsupervised, RNN)



How ML is different?

Data science as different from software
as software was different from hardware.

- Nick Elprin

The rise of Software engineering required inventing new processes like version control, code review, agile, to help teams work effectively. The rise of AI & Machine Learning is now requiring new processes, ..

- Andrew Ng



Key Differences


- ML is metrics driven
- Enormous size of data and models
- Sharing of Models with teams
- Special purpose hardware
- Pipelines
- Difference in best practises



ML is metric driven

- We focus on improving metrics like accuracy, AUC, RUC.
- Aim to make SOTA models, rapidly experiment

ML IS METRICS DRIVEN



Feature-driven Development

- While in software engineering, your team works by building new features.
- Use code-review and agile best practises.



Huge amount of Data dealed

- People work with lots of data for everything.
- Managing datasets are a problem
- Models – a byproduct :P



Huge amount of Data dealed

- People work with TB's of data for everything.
- Managing datasets are a problem
- amount of data generated for a self-driving car
(4 TB/day)



ML models need huge resources

- Most software product take a few minutes to run, which team can easily follow.
- Collaborative challenges faced when working with large teams and remotely.



Data

Schema

Sampling over Time

Volume

+



Model

Algorithms

More Training

Experiments

+



Code

Business Needs

Bug Fixes

Configuration

Pipelines of different stages

- Defining process of steps at each duration like:
 - Input
 - Feature extraction
 - Classification
 - Final output



Special purpose hardware

- Use of special TPUs and cloud Computing
- ML researchers have huge computation needs



Difference in best practises

- Use of Python PEP8 conventions are not always properly practised in jupyter notebooks.
- DataScience = software eng + science&maths

```
from fastai import *
```



Using git for large dataset

- Trying git for dataset of size of 1GB takes about 8-10 minutes to commit
- .gitignore – ignores Dataset versioning



Using git-LFS

“Git Large File Storage (LFS) replaces large files such as audio samples, videos, datasets, and graphics with text pointers inside Git, while storing the file contents on a remote server like GitHub.com or GitHub Enterprise.”



Using git-LFS

- There is a data limit of storing files of 2GB, which is a github limitation.
- Another issue is the ease of placing data files on a cloud storage system (AWS, GCP, etc) as is often required when to run cloud-based AI software.
- Not designed in mind with ML/data science



Other open source tools

- Data Version Control – Track experiments, data
- ML flow – Metric Tracking, Model Deployment
- Neptune.ml – Experiment Tracking, collaboration
- Jovian – Share notebooks, easily reproducible



Why DVC is a good option?

- It's open source
- Has support for multiple cloud providers
- Support for:
 - a) Pipelines
 - b) Versioning ML models
 - c) Versioning Datasets
 - d) Tracking metrics
- “Made for Data scientist, by the Data scientist”



Using DVC with a problem

- We are going to work on a Classic Deep Learning problem like cats vs dogs:

initialising project

```
$ git clone https://github.com/iterative/example-versioning.git
```

```
$ cd example-versioning
```



1. Add the data to dvc

- We use 1000 labelled images of Cats and dogs at first

\$ dvc add data —————> Add datafolder to dvc

\$ python train.py

\$ dvc add model.h5 —————> model is add to dvc



2. Track changes with git

```
$ git add .gitignore model.h5.dvc data.dvc  
metrics.json
```

```
$ git commit -m "First model, trained with 1000  
images"
```

```
$ git tag -a "v1.0" -m "model v1.0, 1000 images"
```



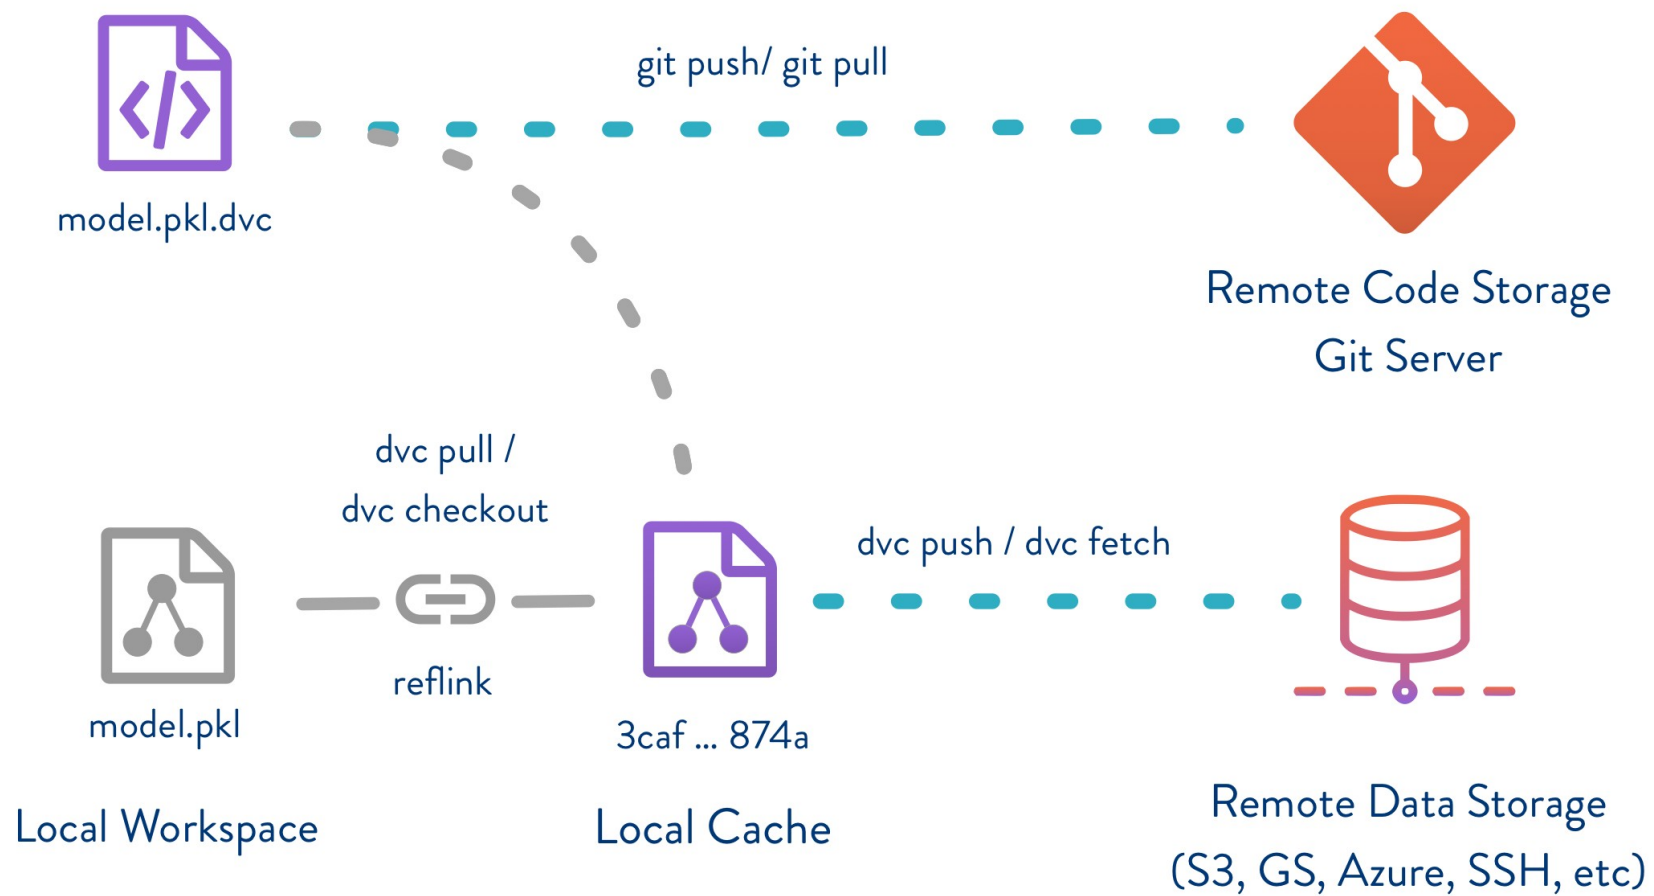
3. Push changes to cloud

\$ dvc push

\$ dvc pull

- Data is pushed to remote server from local cache.
- DVC is storage agnostic and support s3, Azure, GCP, hdfs, ssh





4. Adding more labelled data

(downloading 1000 more images of labelled cat and dog images)

```
$ dvc add more_data
```

```
$ dvc remove model.h5.dvc
```

```
$ python train.py
```

```
$ dvc add model.h5
```



```
$ git add model.h5.dvc data.dvc metrics.json
```

```
$ git commit -m "Second model, trained with  
2000 images"
```

```
$ git tag -a "v2.0" -m "model v2.0, 2000 images"
```




What is Dataset versioning?

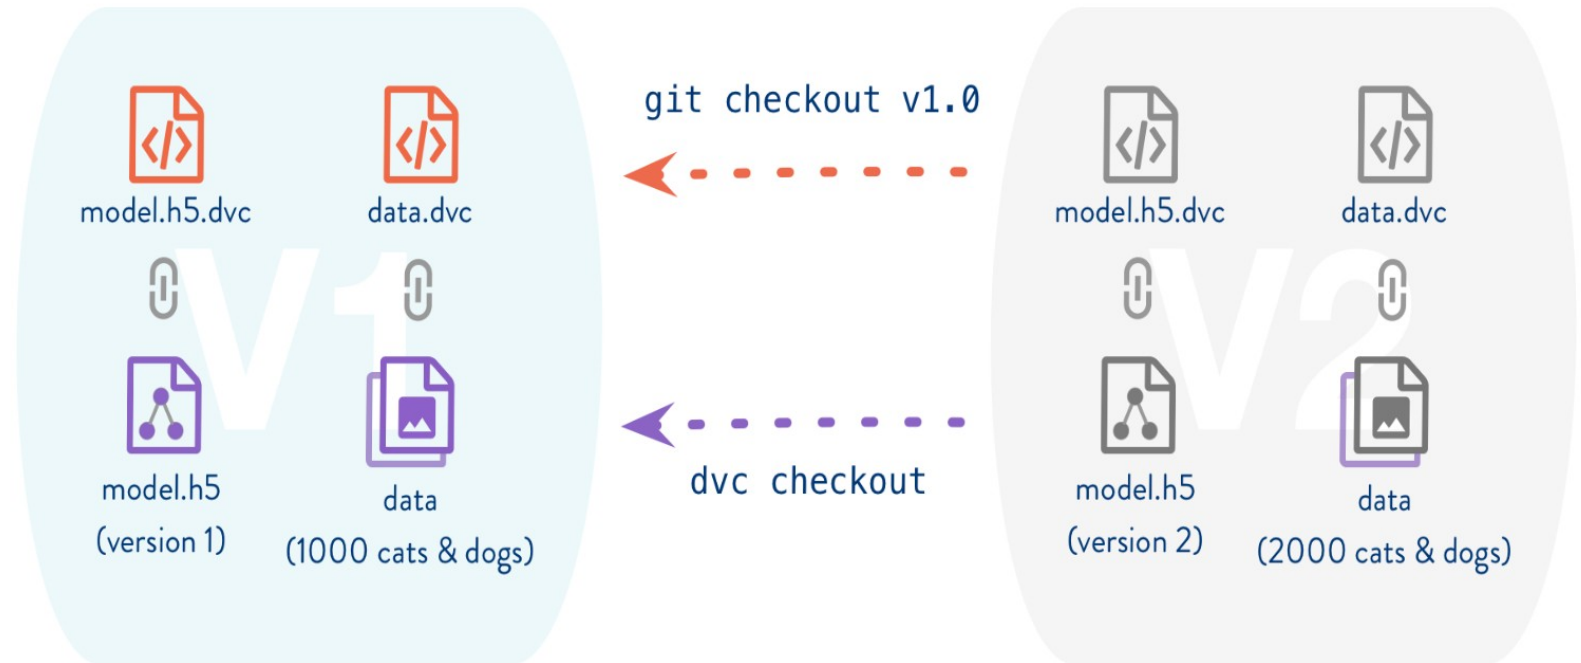
Tracking of data at each time

When working with TB's of data it's essential

Different splits for train/test/val



5. Switching between versions



```
$ git checkout v1.0
```

```
$ dvc checkout
```

6. Running experiments

```
$ dvc pull data_folder.dvc
```

```
$ dvc run <parameters>
```

DVC is internally building a dependency graph using dependencies, output and stores the result.



```
dvc run -f modelDvcfile \  
    -d train.py -d data \  
    -M metrics.json \  
    -o model.h5 \  
-o bottleneck_features_train.npy \  
python train.py
```



d for dependency: specify an input file

o for output: specify an output file ignored by git and tracked by dvc

M for metric: specify an output file tracked by git

f for file: specify the name of the dvc file.



Model versioning

Models = Code + data + hyperparameters

- Everytime our hyperparameters are changed, we can have corresponding DVC files with store the changes
- You can even reproduce the old results



7. More Experiments

- I use VGG architecture for version1
- Then resnet for version2
- V2 made much more improv
Bottle-necks
- Fine tune models



8. Iterating over experiments

- `$ dvc repro train_v2.dvc`
- `$ dvc metrics show`



ML Reproducibility crisis

- Much of ML projects are not reproducible
- Managing the metrics and hyperparameters which changed this module is very important



Conclusion



Questions

